



CHAPTER 7

More SQL: Complex Queries, Triggers, Views, and Schema Modification

Chapter 7 Outline

- More Complex SQL Retrieval Queries
- Specifying Semantic Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Modification in SQL

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping

Comparisons Involving NULL and Three-Valued Logic

- Meanings of `NULL`
 - **Unknown value**
 - **Unavailable or withheld value**
 - **Not applicable attribute**
- Each individual `NULL` value considered to be different from every other `NULL` value
- SQL uses a three-valued logic:
 - TRUE, FALSE, and UNKNOWN (like Maybe)
- **`NULL = NULL` comparison is avoided**

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

Table 7.1 Logical Connectives in Three-Valued Logic

(a)	AND	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	OR	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	NOT			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

Three-Valued Logic

<i>p</i>	<i>q</i>	<i>p OR q</i>	<i>p AND q</i>	<i>p = q</i>
True	True	True	True	True
True	False	True	False	False
True	Unknown	True	Unknown	Unknown
False	True	True	False	False
False	False	False	False	True
False	Unknown	Unknown	False	Unknown
Unknown	True	True	Unknown	Unknown
Unknown	False	Unknown	False	Unknown
Unknown	Unknown	Unknown	Unknown	Unknown

<i>p</i>	NOT p
True	False
False	True
Unknown	Unknown

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- SQL allows queries that check whether an attribute value is NULL
 - IS or IS NOT NULL

Query 18. Retrieve the names of all employees who do not have supervisors.

```
Q18:  SELECT      Fname, Lname
        FROM       EMPLOYEE
        WHERE      Super_ssn IS NULL;
```

Nested Queries, Tuples, and Set/Multiset Comparisons

■ Nested queries

- Complete select-from-where blocks within WHERE clause of another query
- **Outer query and nested subqueries**

■ Comparison operator `IN`

- Compares value v with a set (or multiset) of values V
- Evaluates to `TRUE` if v is one of the elements in V

Nested Queries (cont'd.)

Make a list of all project numbers for projects that involve employee Smith either as worker or as a manager of the department that controls the project:

```
Q4A:  SELECT      DISTINCT Pnumber
      FROM        PROJECT
      WHERE       Pnumber IN
                  ( SELECT      Pnumber
                      FROM        PROJECT, DEPARTMENT, EMPLOYEE
                      WHERE       Dnum=Dnumber AND
                                  Mgr_ssn=Ssn AND Lname='Smith' )
                  OR
                  Pnumber IN
                  ( SELECT      Pno
                      FROM        WORKS_ON, EMPLOYEE
                      WHERE       Essn=Ssn AND Lname='Smith' );
```

Nested Queries (cont'd.)

- Use tuples of values in comparisons
 - Place them within parentheses

```
SELECT      DISTINCT Essn
FROM        WORKS_ON
WHERE       (Pno, Hours) IN ( SELECT      Pno, Hours
                           FROM        WORKS_ON
                           WHERE       Essn='123456789' );
```

Nested Queries (cont'd.)

- Use other comparison operators to compare a single value v
 - $= ANY$ (or $= SOME$) operator [equivalent to IN]
 - Returns TRUE if the value v is equal to some value in the set
 - Other operators that can be combined with ANY (or SOME): $>$, \geq , $<$, \leq , and \neq
 - ALL: value must exceed all values from nested query

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL ( SELECT      Salary
                           FROM        EMPLOYEE
                           WHERE       Dno=5 );
```

General Form of ALL, ANY, SOME

```
SELECT [column_name ]  
FROM [table_name]  
WHERE expression operator  
      {ALL | ANY | SOME} ( subquery )
```

Nested Queries (cont'd.)

- Avoid potential errors and ambiguities
 - Create tuple variables (aliases) for all tables referenced in SQL query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT      E.Fname, E.Lname
        FROM       EMPLOYEE AS E
        WHERE      E.Ssn IN  ( SELECT      Essn
                                FROM       DEPENDENT AS D
                                WHERE      E.Fname=D.Dependent_name
                                            AND E.Sex=D.Sex );
```

Understanding a nested (correlated) query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT      E.Fname, E.Lname
        FROM       EMPLOYEE AS E
        WHERE      E.Ssn IN  ( SELECT      Essn
                                FROM       DEPENDENT AS D
                                WHERE      E.Fname=D.Dependent_name
                                            AND E.Sex=D.Sex );
```

For each E tuple,

Evaluate the nested query

which retrieves the Essn values of all D tuples

with the same sex and name as E tuple

If the Ssn value of E tuple is in the result,

then select the E tuple

Correlated Nested Queries

- **Queries that are nested using the = or IN comparison operator** can be collapsed into one single block: E.g., Q16 can be written as:
- **Q16A:** **SELECT** E.Fname, E.Lname
FROM EMPLOYEE AS E, DEPENDENT AS D
WHERE E.Ssn=D.Essn **AND** E.Sex=D.Sex **AND** E.Fname=D.Dependent_name;
- **Correlated nested query**
 - Evaluated once for each tuple in the outer query

The EXISTS and UNIQUE Functions in SQL for correlating queries

- EXISTS function
 - Check whether the result of a correlated nested query is empty or not. They are Boolean functions that return a TRUE or FALSE result.
- EXISTS and NOT EXISTS
 - Typically used in conjunction with a correlated nested query
- SQL function UNIQUE (Q)
 - Returns TRUE if there are no duplicate tuples in the result of query Q

USE of EXISTS

List the managers who have at least one dependent

Q7:

```
SELECT Fname, Lname  
FROM Employee  
WHERE EXISTS (SELECT *  
                  FROM DEPENDENT  
                  WHERE Ssn= Essn)
```

```
AND EXISTS (SELECT *  
                  FROM Department  
                  WHERE Ssn= Mgr_Ssn)
```

Explicit Sets and Renaming of Attributes in SQL

- Can use explicit set of values in WHERE clause

Q17: **SELECT** **DISTINCT** Essn
 FROM WORKS_ON
 WHERE Pno **IN** (1, 2, 3);

- Use qualifier AS followed by desired new name
 - Rename any attribute that appears in the result of a query

Q8A: **SELECT** E.Lname **AS** Employee_name, S.Lname **AS** Supervisor_name
 FROM EMPLOYEE **AS** E, EMPLOYEE **AS** S
 WHERE E.Super_ssn=S.Ssn;

Specifying Joined Tables in the FROM Clause of SQL

- **Joined table**
 - Permits users to specify a table resulting from a join operation in the FROM clause of a query
- **The FROM clause in Q1A**
 - Contains a single joined table. JOIN may also be called INNER JOIN

Q1A: **SELECT** Fname, Lname, Address
 FROM (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
 WHERE Dname='Research';

Different Types of JOINed Tables in SQL

- Specify different types of join
 - NATURAL JOIN
 - Various types of OUTER JOIN (LEFT, RIGHT, FULL)
- NATURAL JOIN on two relations R and S
 - No join condition specified
 - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

Q1B: **SELECT** Fname, Lname, Address
 FROM (EMPLOYEE NATURAL JOIN
 (DEPARTMENT AS DEPT (Dname, Dno, Mssn,
 Msdate)))
 WHERE Dname='Research';

The above works with EMPLOYEE.Dno = DEPT.Dno as an implicit join condition

INNER and OUTER Joins

- INNER JOIN (**versus** OUTER JOIN)
 - Default type of join in a joined table
 - Tuple is included in the result only if a matching tuple exists in the other relation
- LEFT OUTER JOIN
 - Every tuple in left table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of right table
- RIGHT OUTER JOIN
 - Every tuple in right table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of left table

Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
 - **COUNT, SUM, MAX, MIN, and AVG**
- **Grouping**
 - Create subgroups of tuples before summarizing
- To select entire groups, **HAVING clause** is used
- Aggregate functions can be used in the **SELECT clause** or in a **HAVING clause**

Renaming Results of Aggregation

- Following query returns a single row of computed values from EMPLOYEE table:

Q19: **SELECT** **SUM** (Salary), **MAX** (Salary), **MIN** (Salary), **AVG** (Salary)
 FROM EMPLOYEE;

- The result can be presented with new names:

Q19A: **SELECT** **SUM** (Salary) **AS** Total_Sal, **MAX** (Salary) **AS** Highest_Sal, **MIN** (Salary) **AS** Lowest_Sal, **AVG** (Salary) **AS** Average_Sal
 FROM EMPLOYEE;

Aggregate Functions in SQL (cont'd.)

- NULL values are discarded when aggregate functions are applied to a particular column

Query 20. Find the sum of the salaries of all employees of the ‘Research’ department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:   SELECT      SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
          FROM        (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
          WHERE       Dname='Research';
```

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the ‘Research’ department (Q22).

```
Q21:   SELECT      COUNT (*)
          FROM        EMPLOYEE;
```

```
Q22:   SELECT      COUNT (*)
          FROM        EMPLOYEE, DEPARTMENT
          WHERE       DNO=DNUMBER AND DNAME='Research';
```

Grouping: The GROUP BY Clause

- Partition relation into subsets of tuples
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- GROUP BY clause
 - Specifies grouping attributes
- COUNT (*) counts the number of rows in the group

Examples of GROUP BY

- The grouping attribute must appear in the SELECT clause:

Q24: **SELECT** Dno, **COUNT (*)**, **AVG** (Salary)
 FROM EMPLOYEE
 GROUP BY Dno;

- If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query)
- GROUP BY may be applied to the result of a JOIN:

Q25: **SELECT** Pnumber, Pname, **COUNT (*)**
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname;

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

■ **HAVING clause**

- Provides a condition to select or reject an entire group:
- **Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

Q26: **SELECT** Pnumber, Pname, **COUNT (*)**
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname
 HAVING **COUNT (*) > 2;**

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.
- INCORRECT QUERY:

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary>40000
GROUP BY    Dno
HAVING     COUNT (*) > 5;
```

Combining the WHERE and the HAVING Clause (continued)

Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

Query 28. For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

Q28: **SELECT** Dnumber, COUNT (*)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber=Dno AND Salary>40000 AND
 (**SELECT** Dno
 FROM EMPLOYEE
 GROUP BY Dno
 HAVING COUNT (*) > 5)

Use of CASE

- SQL also has a CASE construct
- Used when a value can be different based on certain conditions.
- Can be used in any part of an SQL query where a value is expected
- Applicable when querying, inserting or updating tuples

EXAMPLE of use of CASE

- The following example shows that employees are receiving different raises in different departments
(A variation of the update U6)

- U6':

UPDATE	EMPLOYEE
SET	Salary =
CASE	WHEN Dno = 5 THEN Salary + 2000
	WHEN Dno = 4 THEN Salary + 1500
	WHEN Dno = 1 THEN Salary + 3000

EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>
FROM <table list>
[ WHERE <condition> ]
[ GROUP BY <grouping attribute(s)> ]
[ HAVING <group condition> ]
[ ORDER BY <attribute list> ];
```

Specifying Constraints as Assertions and Actions as Triggers

- Semantic Constraints: The following are beyond the scope of the EER and relational model
- **CREATE ASSERTION**
 - Specify additional types of constraints outside scope of built-in relational model constraints
- **CREATE TRIGGER**
 - Specify automatic actions that database system will perform when certain events and conditions occur

Specifying General Constraints as Assertions in SQL

■ CREATE ASSERTION

- Specify a query that selects any tuples that violate the desired condition
- Use only in cases where it goes beyond a simple CHECK which applies to individual attributes and domains
- Salary of an employee must be less than the manager

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM        EMPLOYEE E, EMPLOYEE M,
                                  DEPARTMENT D
                     WHERE      E.Salary > M.Salary
                               AND E.Dno = D.Dnumber
                               AND D.Mgr_ssn = M.Ssn ) );
```

Introduction to Triggers in SQL

- CREATE TRIGGER statement
 - Used to monitor the database
- Typical trigger has three components which make it a rule for an “active database” (more on active databases in section 26.1) :
 - **Event(s)**
 - **Condition**
 - **Action**

USE OF TRIGGERS

- AN EXAMPLE with standard Syntax. (Note : other SQL implementations like PostgreSQL use a different syntax.)

R5:

```
CREATE TRIGGER SALARY_VIOLATION  
BEFORE INSERT OR UPDATE OF Salary, Supervisor_ssn ON  
EMPLOYEE
```

FOR EACH ROW

```
WHEN (NEW.Salary > ( SELECT Salary FROM EMPLOYEE  
                      WHERE Ssn = NEW.Supervisor_Ssn))  
INFORM_SUPERVISOR (NEW.Supervisor.Ssn, New.Ssn)
```

Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table derived from other tables called the **defining tables**
 - Considered to be a virtual table that is not necessarily populated

Specification of Views in SQL

■ CREATE VIEW command

- Give table name, list of attribute names, and a query to specify the contents of the view
- In V1, attributes retain the names from base tables. In V2, attributes are assigned names

V1: CREATE VIEW WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn=Essn AND Pno=Pnumber;

V2: CREATE VIEW DEPT_INFO(Dept_name, No_of_emps, Total_sal)
 AS SELECT Dname, COUNT (*), SUM (Salary)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber=Dno
 GROUP BY Dname;

Specification of Views in SQL (cont'd.)

- Once a View is defined, SQL queries can use the View relation in the FROM clause
- View is always up-to-date
 - Responsibility of the DBMS and not the user
- **DROP VIEW** command
 - Dispose of a view

Schema Change Statements in SQL

■ Schema evolution commands

- DBA may want to change the schema while the database is operational
- Does not require recompilation of the database schema

The DROP Command

- **DROP command**
 - Used to drop named schema elements, such as tables, domains, or constraint
- **Drop behavior options:**
 - CASCADE and RESTRICT
- **Example:**
 - `DROP SCHEMA COMPANY CASCADE;`
 - This removes the schema and all its elements including tables, views, constraints, etc.
 - RESTRICT: drops only nothing in it

The ALTER table command

- **Alter table actions** include:
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);`

Adding and Dropping Constraints

- Change constraints specified on a table
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```

Dropping Columns, Default Values

- To drop a column
 - Choose either CASCADE or RESTRICT
 - CASCADE would drop the column from views etc.
 - RESTRICT is possible if no views refer to it.
- Default values can be dropped and altered :
ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN Address CASCADE;
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn DROP DEFAULT;
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn SET DEFAULT '333445555';

Table 7.2 Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]
    { , <column name> <column type> [ <attribute constraint> ] }
    [ <table constraint> { , <table constraint> } ] )
```

```
DROP TABLE <table name>
```

```
ALTER TABLE <table name> ADD <column name> <column type>
```

```
SELECT [ DISTINCT ] <attribute list>
```

```
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }
[ WHERE <condition> ]
```

```
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]
```

```
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) )
    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) ) ) )
```

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]
```

```
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }
| <select statement> )
```

continued on next slide

Table 7.2 (continued)

Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>

[WHERE <selection condition>]

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>

ON <table name> (<column name> [<order>] { , <column name> [<order>] })

[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]

AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

Summary

- Complex SQL:
 - Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- Handling semantic constraints with CREATE ASSERTION and CREATE TRIGGER
- CREATE VIEW statement and materialization strategies
- Schema Modification for the DBAs using ALTER TABLE , ADD and DROP COLUMN, ALTER CONSTRAINT etc .